

SMBC: Engineering a Fast Solver in OCaml

Simon Cruanes

Veridis, Inria Nancy
<https://cedeela.fr/~simon/>

28th of March, 2017

Summary

- 1 Presentation of SMBC ("Satisfiability Modulo Bounded Checking")
- 2 Implementation
- 3 Profiling for Better Performance

Example problem

Given:

- inductive type declarations
 - (recursive) function definitions
 - a **goal**: an expression with variables in it
- find assignment of variables satisfying the goal

Example problem

Given:

- inductive type declarations
 - (recursive) function definitions
 - a **goal**: an expression with variables in it
- find assignment of variables satisfying the goal

Example

Ask the solver to find a palindrome list of length 2 (e.g. [1;1]).

```
let rec length = function
| [] -> 0
| _ :: tail -> succ (length tail)
```

```
let rec rev = function
| [] -> []
| x :: tail -> rev tail @ [x]
```

(* magic happens here *)

goal (rev l = l && length l = 2)

More examples

Example

Ask the solver to find a regex matching "aabb"

```
type char = A | B
type string = char list
type regex =
| Epsilon (* empty *)
| Char of char
| Star of regex
| Or of regex * regex (* choice *)
| Concat of regex * regex (* concatenation *)

let rec match_re : regex -> string -> bool = ...
goal (match_re r [A;A;B;B])
```

We get $r = (\epsilon \mid a^*) \cdot b^*$, i.e.

$r = \text{Concat}(\text{Or}(\text{Epsilon}, \text{Star}(\text{Char } A)), \text{Star}(\text{Char } B))$

More Examples

Example

Solving a sudoku

```
type cell = C1 | C2 | ... | C9
type 'a sudoku = 'a list list

let rec is_instance : cell sudoku -> cell option sudoku -> bool = (* ... *)
let rec is_valid : cell sudoku -> bool = (* ... *)
let partial_sudoku : cell option sudoku = [[None; Some C1; ...]; ...; ]
(* find a full sudoku that matches "partial_sudoku" *)
goal (is_instance e partial_sudoku && is_valid e)
```

More Examples

Example

Solving a sudoku

```
type cell = C1 | C2 | ... | C9
type 'a sudoku = 'a list list

let rec is_instance : cell sudoku -> cell option sudoku -> bool = (* ... *)
let rec is_valid : cell sudoku -> bool = (* ... *)
let partial_sudoku : cell option sudoku = [[None; Some C1; ...]; ...; ]
(* find a full sudoku that matches "partial_sudoku" *)
goal (is_instance e partial_sudoku && is_valid e)
```

- **combinatorial explosion**, large search space
- write a **SMT solver** (satisfiability modulo theory)
- solves in 14 s (not bad for a general-purpose tool)

Similar Tools

HBMC : source of inspiration, bit-blasting Haskell → SAT

small check : native code, tries all values up to depth k

lazy small check : same, but uses lazyness to expand

narrowing : similar to LSC, refine meta-variables on demand

CVC4 : handles datatypes and recursive functions by quantifier instantiation + finite model finding (\rightarrow inefficient?)

QuickCheck & co : random generation of inputs. Very bad on tight constraints.

...

Draw inspiration from HBMC / narrowing+SAT.

The Bigger Picture

- make a better solver for problems based on **recursive functions**
- the rest of the talk: **implementation**
- use **SMT** techniques for not drowning in search space
 - “Satisfiability Modulo Bounded Checking”
- relation to **Nunchaku** (model finder for HO logic):
 - ▶ SMBC is a backend
 - ▶ other backends not very good on this fragment
- useful and widely applicable problem!
- paper submitted to **CADE**

Summary

- 1 Presentation of SMBC ("Satisfiability Modulo Bounded Checking")
- 2 Implementation
- 3 Profiling for Better Performance

Using a SAT-solver Library

Use a **SAT solver**, here .

→ does the backtracking and exploration.

Modularity

MSat is a *library* using an OCaml functor.

Bring your own theory!

```
module type THEORY = sig
  type formula
  type clause = (formula * bool) list

  type result = Ok of clause list | Conflict of clause

  val assume : formula -> bool -> result (* my code *)
end
```

```
module Sat(T:THEORY) : sig
  val solve : formula list -> bool (* library code *)
end
```

- one big Solver module (3,500 loc)
- more than 20 mutually recursive types at the beginning
(to replace tables by embedding data inside objects)
 - yes, *ignoring* the rules of SW engineering can be fine
 - ... for performance reasons
- one can write C in any language, even OCaml!

A few design Decisions

- terms (“expressions”, that is, trees) are DAGs with perfect sharing
 - save memory, constant-time comparison
 - most provers/SMT do it, even in C
- interpreter with **caching** of normal form of a value
 - caching often dramatically improves performance (when it applies)
- **backtracking**: a big stack of “undo” functions
 - low memory footprint, low overhead

Summary

- 1 Presentation of SMBC ("Satisfiability Modulo Bounded Checking")
- 2 Implementation
- 3 Profiling for Better Performance

How to debug Performance

SMBC needs to be very efficient, because it is *kind of* bruteforce.

Rules of thumb

- try to use efficient algorithms everywhere
- try to avoid allocating too much
- use compiler optimizations (here, ocaml+flambda)
- avoid obviously inefficient code, **but**:
- avoid “premature optimization” (as would say D. Knuth)
→ hence the need for profiling

CPU profiling: “perf”

perf: standard tool on Linux (initially there for the kernel)

CPU profiling: “perf”

perf: standard tool on Linux (initially there for the kernel)

```
% perf record --call-graph=dwarf ./smbc.native examples/ty_infer.smt2
(result SAT
:model ((val e_2
  (lam
    (lam
      (lam
        (app (lam (app (var (s (s z)))) (var z) b))
          (app (var (s z)) (var z) a) b)))))))
[ perf record: Woken up 40 times to write data ]
[ perf record: Captured and wrote 9.946 MB perf.data (1234 samples) ]

% perf report
```

Perf (cont'd)

```
Samples: 1K of event 'cycles:u', Event count (approx.): 1069569901
      Children   Self  Command  Shared Object    Symbol
+ 99.12%  0.00% smbc.native  smbc.native  [.]. caml_main
+ 99.12%  0.00% smbc.native  smbc.native  [.]. main
+ 99.12%  0.00% smbc.native  libc-2.24.so  [.]. __libc_start_main
+ 99.12%  0.00% smbc.native  smbc.native  [.]. _start
+ 99.05%  0.00% smbc.native  smbc.native  [.]. camlSmbc__entry
+ 99.05%  0.00% smbc.native  smbc.native  [.]. caml_program
+ 99.05%  0.00% smbc.native  smbc.native  [.]. caml_start_program
+ 98.98%  0.00% smbc.native  smbc.native  [.]. camlSmbc__solve_1377
+ 98.94%  0.04% smbc.native  smbc.native  [.]. camlSolver__iter_7227
+ 98.84%  0.00% smbc.native  smbc.native  [.]. camlMsat__External__solve_inner_4075
- 98.75%  1.50% smbc.native  smbc.native  [.]. camlMsat__Internal__search_1827
- 97.25% camlMsat__Internal__search_1827
  + 46.61% camlMsat__Internal__theory_propagate_1801
  + 21.72% camlMsat__Internal__pick_branch_lit_1815
  + 13.89% camlMsat__Internal__propagate_atom_1763
  + 13.46% camlMsat__Internal__propagate_1802
  + 1.49% camlMsat__Internal__add_boolean_conflict_1716
+ 1.50% _start
+ 46.61%  0.25% smbc.native  smbc.native  [.]. camlMsat__Internal__theory_propagate_1801
+ 45.95%  0.08% smbc.native  smbc.native  [.]. camlSolver__assume_5801
+ 41.55%  0.08% smbc.native  smbc.native  [.]. camlList__iter_1252
+ 40.98%  0.16% smbc.native  smbc.native  [.]. camlSolver__update_5637
+ 35.51%  0.58% smbc.native  smbc.native  [.]. camlSolver__compute_nf_add_5401
+ 35.43%  1.90% smbc.native  smbc.native  [.]. camlSolver__compute_nf_noncached_5399
+ 34.79%  2.41% smbc.native  smbc.native  [.]. camlSolver__compute_builtin_5402
+ 21.72%  2.40% smbc.native  smbc.native  [.]. camlMsat__Internal__pick_branch_lit_1815
- 18.91%  4.39% smbc.native  smbc.native  [.]. camlMsat__Iheap__remove_min_1307
- 14.53% camlMsat__Iheap__remove_min_1307
  + 10.93% camlMsat__Internal__f_weight_1510
  + 2.12% camlPervasives__max_1030
  1.07% caml_modify
  + 4.39% _start
+ 17.92%  0.00% smbc.native  smbc.native  [.]. camlSolver__compute_nf_app_5400
```

Perf and Flamegraphs


With deep recursive calls, perf report isn't very good.

Perf and Flamegraphs

With deep recursive calls, perf report isn't very good.

→ flame graphs (<http://www.brendangregg.com/flamegraphs.html>)

```
perf script \
| stackcollapse-perf --kernel \
| sed 's/caml//g' \
| flamegraph > perf.svg
```



Memory Profiling

OCaml has a GC, so I need to minimize allocations.

→ use **spacetime**, a new memory profiler!

- <https://caml.inria.fr/pub/docs/manual-ocaml/spacetime.html>
- <https://blogs.janestreet.com/a-brief-trip-through-spacetime/>

Memory Profiling

OCaml has a GC, so I need to minimize allocations.

→ use **spacetime**, a new memory profiler!

- <https://caml.inria.fr/pub/docs/manual-ocaml/spacetime.html>
- <https://blogs.janestreet.com/a-brief-trip-through-spacetime/>

Available as an OCaml *compiler switch*


```
% opam sw 4.04.0+spacetime  
  
% make clean all  
  
% OCAML_SPACETIME_INTERVAL=100 ./smbc.native examples/ty_infer.smt2  
  
% prof_spacetime serve spacetime-<PID> -e smbca.native
```

Memory Profiling (cont'd)

smbc.native

Mouse over the graph to show where values were allocated. Values allocated from non-OCaml code have their mouse-over popup text in green. Click a portion of the graph to move up the stack.

[Live words](#) [Live blocks](#) [All allocated words](#)



Backtrace (oldest frame first):

- [\(top of stack\)](#)

Horizontal: time

Vertical: space

Colors: track memory allocated from a given program position

Memory Profiling (cont'd)


smbc.native

Mouse over the graph to show where values were allocated. Values allocated from non-OCaml code have their mouse-over popup text in green. Click a portion of the graph to move up the stack.

Live words

Live blocks

All allocated words



Backtrace (oldest frame first):

- [set.ml:364,41--67 \(camISet_elements_aux_1474\)](#)
- [\(top of stack\)](#)

Can zoom into any region!

Memory Profiling

Found a performance bug this way:

- MSat uses a lot of dynamic arrays
- mistake in criterion for re-sizing
 - resize at every Vec.push!
- almost all allocations came from there;
spacetime made it obvious.

Conclusion

- good algorithms (SAT solver here) trump excellent implementation
... but implementation still important!
- OCaml can have reasonable performance **if** used properly
- profile before micro-optimizing
- tooling for profiling is tremendously useful

`perf`: can be used with many languages, de-facto standard on Linux

`spacetime`: awesome, but limited to OCaml

`others`: can also profile by manually inserting counters